

ECE/CS 559: Neural Networks Homework 4


Samuele Pasquale

September 24, 2024

1. Question 1:

(a) The code and the plots for the digits from 0 to 9 are shown below.

```
1 def print_example_each_digit(Xraw, Yraw):
2     """
3     Function that shows one example of all the digits from 0 to 9
4     :param Xraw: Xraw matrix
5     :param Yraw: Yraw matrix
6     :return:
7     """
8     # found flag
9     found = 0
10
11     # search up to digit 9
12     while found < 10:
13         # scan Yraw
14         for i in range(len(Yraw)):
15             # check if it is the target digit
16             if Yraw[i] == found:
17                 # print the image
18                 print_image(Xraw[i], found)
19                 # update flag
20                 found += 1
21                 # exiting the for loop
22                 break
23
24
25 def print_image(row, number):
26     """
27     Function that shows the image of a given value
28     :param row: row of Xraw, with 28x28 pixels
29     :param number: Actual number shown in the image
30     :return:
31     """
32     image = row.reshape(28, 28) # MNIST images are 28x28 pixels
33     plt.imshow(image, cmap='gray')
34     plt.title(f"Label: {number}")
35     plt.show()
```



(b) The code is shown below.

```

1      # set d to 10
2      d = 10
3
4      # generation of M, X, and Y
5      M = generate_M(d, 784, 255*d)
6      X = generate_X(M, Xraw)
7      Y = generate_Y(10, Yraw)
8
9      # Moore-Penrose pseudo inverse computation
10     W = MoorePensorePseudoInverse(X, Y)

```

```

1 def generate_M(rows, cols, divisor, uniform_lower_bound=0, uniform_upper_bound
  =1):
2     """
3     Function that generates a rowsXcols matrix filled of random uniform values
4     between the two
5     given boundaries.
6     :param rows: number of rows
7     :param cols: numeber of columns
8     :param divisor: divisor for M/M
9     :param uniform_lower_bound: lower bound of uniform function
10    :param uniform_upper_bound: upper bound of uniform function
11    :return: return the generated matrix M
12    """
13    # generation of random matrix M
14    M = np.random.uniform(uniform_lower_bound, uniform_upper_bound, (rows, cols)
15    )
16    # return matrix M
17    return M / divisor
18
19 def generate_X(M, Xraw):
20     """
21     Function that generates X matrix as M * Xraw^T
22     :param M: M matrix
23     :param Xraw: Xraw matrix
24     :return: matrix with dimension d * 70,000
25     """

```

```

25     # return  M * Xraw^T
26     return M @ Xraw.T
27
28
29 def generate_Y(rows, Yraw):
30     """
31     Function that generates Y matrix, where each i-column represents
32     the one-hot encoding of each i-row of Yraw
33     :param rows: number of rows
34     :param Yraw: Yraw matrix
35     :return: Y matrix
36     """
37     # Y matrix zeroed out
38     Y = np.zeros((rows, Yraw.shape[0]))
39     # one-hot encoding
40     Y[Yraw, np.arange(Yraw.shape[0])] = 1
41     # return matrix Y
42     return Y
43
44
45 def MoorePensorePseudoInverse(X, Y):
46     """
47     Function that computes the moore pensore pseudo-inverse given X and Y
48     matrices
49     :param X: X matrix
50     :param Y: Y matrix
51     :return: W = Y X^T (X X^T)^(-1)
52     """
53     # pseudo inverse with linalg library
54     X_pseudo_inv = np.linalg.pinv(X)
55     # weights computation
56     W = Y @ X_pseudo_inv
57     # return weights
58     return W

```

- (c) The number of errors and the MSE of the predictor are reported below. As can be noticed from the table, increasing the parameter d causes the number of errors to decrease up to a certain value, 200, beyond which the number of errors saturates at around 10,000. Therefore, increasing d beyond 200 doesn't result in a significant improvement in the number of errors. In conclusion, $d=200$ is a good choice, while a higher value doesn't further decrease the error rate of the predictor.

Finally, it can be noted that the error probability, considering 36,522 errors out of 70,000, is 52.17 %. On the other hand, if the classification were to occur randomly, the error probability would be 9 out of 10 for each image, resulting in a probability of 90 %. Therefore, it can be concluded that the obtained model is still more accurate than a random choice.

d	MSE	Number of Errors
10	0.7645	36,522
50	0.5372	15,856
100	0.4649	12,467
200	0.4221	10,901
500	0.3940	10,399

- (d) In the table 1, the number of errors is shown. As can be noticed the number of errors is higher than the number of errors obtained in (c) with $d = 100$.

The approach taken to improve the accuracy of the algorithm involved training on different numbers of epochs and for different η , with the results shown in Figure 6 and Table 2. From the results obtained, it is concluded that a smaller η allows for slower convergence with fewer oscillations towards a local minimum. However, as shown in the graphs, while the MSE for $\eta = 0.001$ and $\eta = 0.005$ are similar, if η is reduced too much, the updates during training will be too small, and thus for a fixed number of epochs, the model might not escape one minimum to reach a better one. Therefore, in such cases, a larger η would allow for larger steps towards obtaining the optimal solution.

Finally, the impact of the number of epochs on the final result is analyzed. It is observed that, given an optimal η , as described previously, the number of epochs has relatively little impact, and for all tests, the MSE remains of the same order of magnitude. Thus, an excessive number of epochs not only affects the execution time, making it unnecessary for the learning phase, but could even cause an error known as overfitting (source: <https://medium.com/>). In this case, the model becomes excessively adapted to the training data, subsequently losing accuracy on the test data.

d	Number of Errors
100	13918

Table 1: Number of errors for different d values

Errors	Epochs	Learning Rate (eta)
25524	10	0.0001
17802	10	0.001
14541	10	0.005
22990	20	0.0001
15670	20	0.001
12860	20	0.005
20260	30	0.0001
15262	30	0.001
12973	30	0.005
20404	40	0.0001
14082	40	0.001
12271	40	0.005
18896	50	0.0001
13783	50	0.001
12465	50	0.005

Table 2: Comparison of Errors, Epochs, and Learning Rate (eta)




Figure 1: 10 epochs

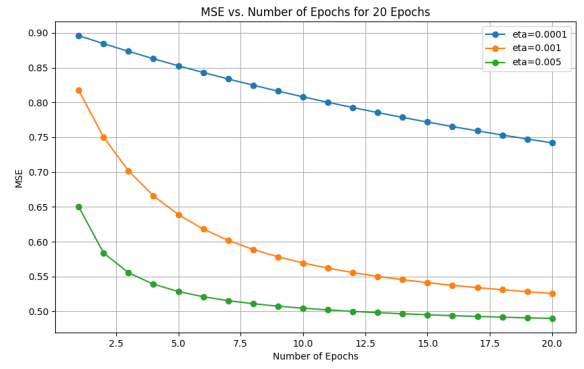


Figure 2: 20 epochs

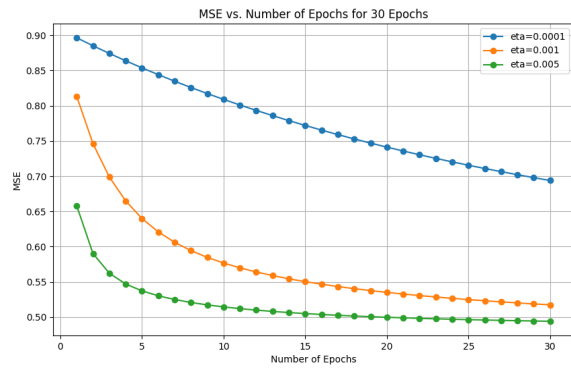


Figure 3: 30 epochs




Figure 4: 40 epochs

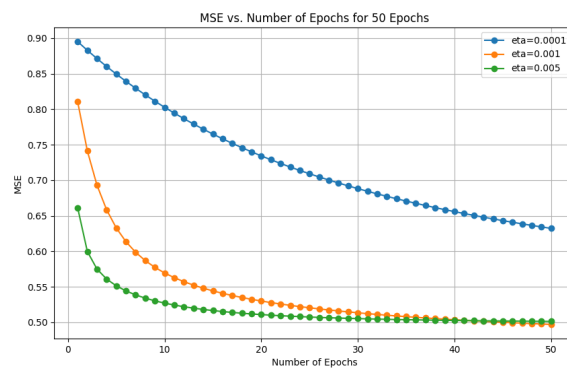


Figure 5: 50 epochs

Figure 6: Tentative of algorithm improvements

